



COMP2004 Distributed Systems

Soumyabrata DEV


School of Computer Science
University College Dublin

<https://soumyabrata.dev/>

University College Dublin (UCD)
4-May-2021

Course website is available here:

<https://soumyabrata.dev/distributedsystems.html>



[Home](#)

COMP2004 Distributed Systems

Date: 4-May-2021

Time: 10:00 hours

Duration: 8 minutes (approximately)

Venue: Online via zoom.


[Handout](#)

[Course Slides](#)

(1/3) Distributed system

Distributed System

A distributed system is a computing environment in which various components are spread across multiple computers on a network¹.




¹Image archived from https://www.splunk.com/en_us/data-insider/what-are-distributed-systems.html


What is time?

Time is a source of order - it allows us to define the order of operations.

Global clock of total order



Local clock with skew and drift



(2/3) Lamport's logical clock

Lamport's "happened-before" notation

$a \rightarrow b$ event a happened before event b

eg. a : message being sent, b : message receipt

Transitive:

if $a \rightarrow b$ and $b \rightarrow c$, then $a \rightarrow c$

Assign “clock” value to each event

- if $a \rightarrow b$ then $clock(a) < clock(b)$
- since time cannot run backwards

If a and b occur on different processes that do not exchange messages, then neither $a \rightarrow b$ nor $b \rightarrow a$ are true.

- These events are called **concurrent**.

Lamport's clocks are defined as:

1. Events within a process is assigned an unique ID.
2. When sending a message from A , we set $clock(A) = clock(A) + 1$, then pass it as a part of the message.
3. On receipt of message in B , we set $clock(B) = \max[clock(B), clock(A)] + 1$.

(3/3) Illustration of Lamport's logical clock

Example

increments by 6

P0



increments by 8

P1



increments by 10

P2



Example

increments by 6

P0

0
6

increments by 8

P1

0
8

increments by 10

P2

0
10

P0 sends message A to P1

Example

increments by 6

P0




increments by 8

P1



A



increments by 10

P2



P1 receives message A (everything is OK since $6 < 16$)

Example

increments by 6

P0

0
6
12
18

increments by 8

P1

0
8
16
24

increments by 10

P2

0
10
20
30

P1 sends message B to P2

Example

increments by 6

P0

0
6
12
18
24

increments by 8

P1

0
8
16
24
30

increments by 10

P2

0
10
20
30
40

B

P2 receives message B (everything is OK since $24 < 40$)

Example

increments by 6

P0

0
6
12
18
24
30

increments by 8

P1

0
8
16
24
32
40

increments by 10

P2

0
10
20
30
40
50

Example

increments by 6

P0

0
6
12
18
24
30
36

increments by 8

P1

0
8
16
24
30
40
48

increments by 10

P2

0
10
20
30
40
50
60


P2 sends message C to P1

Example

increments by 6

increments by 8

increments by 10




P1 receives message C (Ouch! The message was sent at time 60 but received at time 56)

Example

increments by 6

increments by 8

increments by 10



Logical Time at P1 updated to be 1 greater than the time C was sent at.

Example

increments by 6

P0

12
18
24
30
36
42
48

increments by 8

P1

16
24
30
40
48
61
69

increments by 10

P2

20
30
40
50
60
70
80


P1 sends message D to P0

Example

increments by 6

increments by 8

increments by 10




P0 receives message D (Ouch! $56 < 69$)

Example

increments by 6

increments by 8

increments by 10



Logical Time at P0 updated

Example

increments by 6

P0

24
30
36
42
48
70
76

increments by 8

P1

30
40
48
61
69
77
85

increments by 10

P2

40
50
60
70
80
90
100

End of Run

- Lamport clocks can guarantee that if $a \rightarrow b$ then $clock(a) < clock(b)$.
 - However it can't guarantee, that if $clock(a) < clock(b)$ then event a happened before b .
- In order to overcome this limitation, we use **vector clock** to represent the causality of events.



Thank You