



COMP2004 Distributed Systems

Soumyabrata DEV


School of Computer Science
University College Dublin

<https://soumyabrata.dev/>

University College Dublin (UCD)
4-May-2021

Course website is available here:

<https://soumyabrata.dev/distributedsystems.html>



[Home](#)

COMP2004 Distributed Systems

Date: 4-May-2021

Time: 10:00 hours

Duration: 8 minutes (approximately)

Venue: Online via zoom.


[Handout](#)

[Course Slides](#)

(1/3) Distributed system

Distributed System

A distributed system is a computing environment in which various components are spread across multiple computers on a network¹.




¹Image archived from https://www.splunk.com/en_us/data-insider/what-are-distributed-systems.html


What is time?

Time is a source of order - it allows us to define the order of operations.

Global clock of total order



Local clock with skew and drift



(2/3) Lamport's logical clock

Lamport's "happened-before" notation

$a \rightarrow b$ event a happened before event b

eg. a : message being sent, b : message receipt

Transitive:

if $a \rightarrow b$ and $b \rightarrow c$, then $a \rightarrow c$

Assign “clock” value to each event

- if $a \rightarrow b$ then $clock(a) < clock(b)$
- since time cannot run backwards

If a and b occur on different processes that do not exchange messages, then neither $a \rightarrow b$ nor $b \rightarrow a$ are true


- These events are called **concurrent**

Lamport's clocks are defined as:


1. Events within a process is assigned an unique ID. It is a counter that increments from 1.
2. When sending a message from A , we set $clock(A) = clock(A) + 1$, then pass it as a part of the message.
3. On receipt of message in B , we set $clock(B) = \max(clock(B)+1, clock(A))$

(3/3) Illustration of Lamport's logical clock

Example




Example




P0 sends message A to P1

Example




P1 receives message A (everything is OK since $6 < 16$)

Example



P1 sends message B to P2

Example



P2 receives message B (everything is OK since $24 < 40$)

Example

P0

0
6
12
18
24
30

P1

0
8
16
24
32
40

P2


0
10
20
30
40
50

Example

P0	P1	P2
0	0	0
6	8	10
12	16	20
18	24	30
24	30	40
30	40	50
36	48	60


P2 sends message C to P1

Example



P1 receives message C (Ouch! The message was sent at time 60 but received at time 56)

Example




Logical Time at P1 updated to be 1 greater than the time C was sent at.

Example

P0	P1	P2
12	16	20
18	24	30
24	30	40
30	40	50
36	48	60
42	61	70
48	69	80


P1 sends message D to P0

Example



P0 receives message D (Ouch! $56 < 69$)

Example



Logical Time at P0 updated

Example

P0

24
30
36
42
48
70
76

P1

30
40
48
61
69
77
85

P2

40
50
60
70
80
90
100

End of Run

- Lamport clocks can guarantee that if $a \rightarrow b$ then $C(a) < C(b)$.
 - However it can't guarantee, that if $C(a) < C(b)$ then event a happened before b .
- In order to overcome this limitation, we use **vector clock** to represent the causality of events.



Thank You